

Problem 1606: Find Servers That Handled Most Number of Requests

Problem Information

Difficulty: Hard

Acceptance Rate: 45.02%

Paid Only: No

Tags: Array, Heap (Priority Queue), Simulation, Ordered Set

Problem Description

You have `k` servers numbered from `0` to `k-1` that are being used to handle multiple requests simultaneously. Each server has infinite computational capacity but **cannot handle more than one request at a time**. The requests are assigned to servers according to a specific algorithm:

* The `ith` (0-indexed) request arrives. * If all servers are busy, the request is dropped (not handled at all). * If the `(i % k)`th server is available, assign the request to that server. * Otherwise, assign the request to the next available server (wrapping around the list of servers and starting from 0 if necessary). For example, if the `ith` server is busy, try to assign the request to the `(i+1)th` server, then the `(i+2)th` server, and so on.

You are given a **strictly increasing** array `arrival` of positive integers, where `arrival[i]` represents the arrival time of the `ith` request, and another array `load`, where `load[i]` represents the load of the `ith` request (the time it takes to complete). Your goal is to find the **busiest server(s)**. A server is considered **busiest** if it handled the most number of requests successfully among all the servers.

Return _a list containing the IDs (0-indexed) of the**busiest server(s)**_. You may return the IDs in any order.

Example 1:

Input: k = 3, arrival = [1,2,3,4,5], load = [5,2,3,3,3] **Output:** [1] **Explanation:** All of the servers start out available. The first 3 requests are handled by the first 3 servers in order. Request 3 comes in. Server 0 is busy, so it's assigned to the next available server, which is 1. Request 4 comes in. It cannot be handled since all servers are busy, so it is dropped. Servers 0 and 2 handled one request each, while server 1 handled two requests. Hence server 1 is the busiest server.

Example 2:

Input: k = 3, arrival = [1,2,3,4], load = [1,2,1,2] **Output:** [0] **Explanation:** The first 3 requests are handled by first 3 servers. Request 3 comes in. It is handled by server 0 since the server is available. Server 0 handled two requests, while servers 1 and 2 handled one request each. Hence server 0 is the busiest server.

Example 3:

Input: k = 3, arrival = [1,2,3], load = [10,12,11] **Output:** [0,1,2] **Explanation:** Each server handles a single request, so they are all considered the busiest.

Constraints:

* `1 <= k <= 105` * `1 <= arrival.length, load.length <= 105` * `arrival.length == load.length` * `1 <= arrival[i], load[i] <= 109` * `arrival` is **strictly increasing**.

Code Snippets

C++:

```
class Solution {
public:
    vector<int> busiestServers(int k, vector<int>& arrival, vector<int>& load) {
        ...
    };
}
```

Java:

```
class Solution {
    public List<Integer> busiestServers(int k, int[] arrival, int[] load) {
```

```
    }  
    }
```

Python3:

```
class Solution:  
    def busiestServers(self, k: int, arrival: List[int], load: List[int]) ->  
        List[int]:
```